

Reading: Chapter 07 of [Chiusano and Bjarnason 2014]

Purely Functional Parallelism


Advanced Programming

Oct 31, 2019

Andrzej Wasowski & Zhoulai Fu

IT University of Copenhagen

Why parallelism



Because of:

- (1) significantly more computational tasks
- (2) Significantly more CPU

Why functional parallelism

- Many parallelism models involve the using of a shared variable, and side effects
- And side effects can be hard to reason, **easy to get wrong**

blue thread	other threads
<code>x = ...;</code> <code>done = true;</code>	<code>while (!done) {}</code> <code>... = x;</code>

The code works reliably with a dumb compiler, but can cause deadlock for many modern compilers!

Reason: Modern compilers compile “While (! done){}” to `tmp=done; while (!tmp){}`

Clarification 1: concurrency vs parallelism:

- ❖ A minor note about concurrency vs parallelism:
 - ❖ Concurrency describes a *problem* — that things need to happen together
 - ❖ Parallelism describes a *solution* — that is based on multiple threads / CPUs
- ❖ Today's class is about *designing* API for parallelism, under the functional paradigm


Clarification2: Parallelism is not always applicable

- ❖ You can calculate $1+2+3+\dots+N$ via parallelism
- ❖ But, you may not calculate $0.1+0.2+0.3 \dots$ via parallelism

Demo

Background: Java's ExecutorService and Future API

```
class ExecutorService {  
    def submit[A] (a: Callable[A]): Future[A]  
}  
trait Future[A] {  
    def get: A  
}
```



https://www.slideshare.net/afkham_azeez/java-colombo-developing-highly-scalable-apps

```
public interface ExecutorService  
extends Executor
```


An **Executor** that provides methods to manage termination and methods that can produce a **Future** for tracking progress of one or more asynchronous tasks.

The thread pool execution uses a blocking queue. It keeps storing all the tasks that you have submitted to the **executor service**, and all threads (workers) are always running and performing the same steps:

- Take the Task from the queue
- Execute it
- Take the next or wait until a task will be added to the queue

Typical usage of ExecutorService

```
class ExecutorService {  
    def submit[A](a: Callable[A]): Future[A]  
}  
trait Future[A] {  
    def get: A  
}
```



Typical usage1

```
val service: ExecutorService =  
    Executors.newFixedThreadPool(2)  
service.submit(t1)  
service.submit(t2)  
service.submit(t3)
```

Typical usage2

```
val es=Executors.newWorkStealingPool()  
val fut=es.submit(new Callable[Int] {  
    override def call(): Int = (1 to 10).sum  
})  
val res=fut.get
```

Background: Strict and Lazy Functions

- ❖ Strict Evaluation (by-value): function argument evaluated before entering the function
- ❖ Lazy evaluation, (by-name, by-need)

`val x = {println ("eager"); 5}`

Eager

`lazy val y = {println ("lazy"); 4}`

Lazy

What would be $x+x+y+y$?

DEMO

Example of using a lazy function (recall)

```
def time[A](a: => A) = {  
  val now = System.nanoTime  
  val result = a  
  val micros = (System.nanoTime - now) / 1000  
  println("%d microseconds".format(micros))  
  result  
}
```



API for functional parallelism

- No right answers in design
- You will see a collection of design choices
- You are to understand their trade-offs, and think critically.

Why not use Java Thread

```
trait Runnable { def run: Unit }
```

```
class Thread(r: Runnable) {  
  def start: Unit  
  def join: Unit  
}
```

Begins running x in
a separate thread.

Blocks the calling thread
until x finishes running.

- ❖ Side-effects are evil when it comes to functional program and reasoning: if we want to get any information out of a `Runnable`, it has to have some side effect, like mutating some state that we can inspect.

Design Goals

- ❖ No right answers in design.
- ❖ **Pure:** function return the same value for the same input, without observable side effects
- ❖ **High-level:** having the capability to write something like foldleft, as in sequential programs.
- ❖

```
def sum(ints: Seq[Int]): Int =  
  ints.foldLeft(0) ((a,b) => a + b)
```

Design Methodologies

- ❖ Start from a very simple use case
- ❖ Try - Challenge - Refine

Example: Summing a list with divide-and-conquer

```
def sum(ints: IndexedSeq[Int]): Int =  
  if (ints.size <= 1)  
    ints.headOption getOrElse 0  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    sum(l) + sum(r)  
  }
```

Divides the sequence in half using the `splitAt` function.

IndexedSeq is a superclass of random-access sequences like `Vector` in the standard library. Unlike lists, these sequences provide an efficient `splitAt` method for dividing them into two parts at a particular index.

headOption is a method defined on all collections in Scala. We saw this function in chapter 4.

Recursively sums both halves and adds the results together.

❖ Listing 7.1, [Chiusano et al]

The Making of a Parallel Sum (1 - try)

```
def sum(ints: IndexedSeq[Int]): Int =  
  if (ints.size <= 1)  
    ints headOption getOrElse 0  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    val sumL: Par[Int] = Par.unit(sum(l))  
    val sumR: Par[Int] = Par.unit(sum(r))  
    Par.get(sumL) + Par.get(sumR)  
  }
```

Computes the left half in parallel.

Computes the right half in parallel.

Extracts both results and sums them.

- ❖ Need a data type to contain parallel computation results: **Par[A]**
- ❖ Need a function to evaluate a computation in a separate thread
 - ❖ **Par.unit** (a: =>A): Par[A]
- ❖ Need another function to extract a result from a Par[A]:
 - ❖ **Par.get** [A] (a: Par[A]):A

The Making of a Parallel Sum (1 - problem)

```
def sum(ints: IndexedSeq[Int]): Int =  
  if (ints.size <= 1)  
    ints headOption getOrElse 0  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    val sumL: Par[Int] = Par.unit(sum(l))  
    val sumR: Par[Int] = Par.unit(sum(r))  
    Par.get(sumL) + Par.get(sumR)  
  }
```

Computes the left half in parallel.

Computes the right half in parallel.

Extracts both results and sums them.

- ❖ For the sake of parallelization, `Par.unit` has to delay the computation until `Par.get`
- ❖ Problem: The whole computation is still sequential because “+” is strict

The Making of a Parallel Sum: (2 - try)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.size <= 1)  
    Par.unit(ints.headOption.getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(sum(l), sum(r))(_ + _)  
  }
```

- ❖ Par.map2 is a new higher-order function for combining the result of two parallel computations.
- ❖ Q: What is its signature?
- ❖ A: `Par.map2[A,B,C] (a: Par[A], b: Par[B]) (f: (A,B) => C): Par[C]`
- ❖ Q: Should Par.map2 be lazy or strict?
- ❖ A: :If it is strict, we'll strictly construct the entire left half of the tree of summations first before moving on to (strictly) constructing the right half ==> Let Par.map2 be lazy

The Making of a Parallel Sum: (2 - problem)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.size <= 1)  
    Par.unit(ints.headOption getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(sum(l), sum(r))(_ + _)  
  }
```

- ❖ Q: Do we always want to evaluate the two arguments to Par.map2 in parallel?
- ❖ A: : Probably not. Consider Par.map2(Par.unit(1), Par.unit(2))(_+_). The overhead for thread creation/management is swamping any tiny gains from parallelization.
- ❖ Problem: This API is ver inexplicit about when computations gets forked off the main thread — the programmer cannot specify where this forking should occur.

The Making of a Parallel Sum: (3 - try)

```
def sum(ints: IndexedSeq[Int]): Par[Int] =  
  if (ints.length <= 1)  
    Par.unit(ints.headOption getOrElse 0)  
  else {  
    val (l,r) = ints.splitAt(ints.length/2)  
    Par.map2(Par.fork(sum(l)), Par.fork(sum(r))) (_ + _)  
  }
```

- ❖ `Par.fork[A](a: => Par[A]): Par[Int]` runs *a* in a separate logical thread
- ❖ With *Par.fork*, we can make *Par.map2* strict, leaving it up to the programmer to wrap arguments if they want

The Making of a Parallel Sum: (final)

- ❖ We let *fork* hold on to its unevaluated argument until later. It takes an unevaluated $Par[A]$ and marks it for concurrent evaluation later
- ❖ In this model, $Par[A]$ holds a *description* of a parallel computation that gets *interpreted* at a later time by something like the *get* function

Implementation: reused API from Java

```
class ExecutorService {  
  def submit[A](a: Callable[A]): Future[A]  
}  
trait Callable[A] { def call: A }  
trait Future[A] {  
  def get: A  
  def get(timeout: Long, unit: TimeUnit): A  
  def cancel(evenIfRunning: Boolean): Boolean  
  def isDone: Boolean  
  def isCancelled: Boolean  
}
```

- ❖ *Future* is a handle for running a computation in a separate thread
- ❖ *ExecutorService* allows us submit a *Callable* value and get back a corresponding *Future*
- ❖ When *Future* obtain a value from *get*, it blocks the current thread until the value is available.
- ❖ *Future* has extra features for cancellation, e.g., throwing an exception after blocking for a certain amount of time.

Implementation: Other API

- ❖ Type alias: *type Par[A] = ExecutorService => Future[A]*
- ❖ Object Par that holds three primitive operations: *unit*, *map2*, and *fork*

```

type Par[A] = ExecutorService => Future[A]

object Par {
  def unit[A](a:A):Par[A] = (es:ExecutorService)=> UnitFuture(a)

  private case class UnitFuture[A] (get: A) extends Future[A] {
    override def cancel(mayInterruptIfRunning: Boolean): Boolean = false

    override def isCancelled: Boolean = false

    override def isDone: Boolean = true

    override def get(timeout: Long, unit: TimeUnit): A = get
  }

  def map2[A,B,C] (a:Par[A],b:Par[B])(f:(A,B)=>C):Par[C] =
    (es:concurrent.ExecutorService) =>{
      val af = a(es)
      val bf = b(es)
      UnitFuture(f(af.get,bf.get))
    }

  def fork[A](a: => Par[A]) : Par[A] =
    es=>es.submit(new Callable[A] {
      override def call = a(es).get
    })
}

```

- unit is represented as a function that returns a UnitFuture, which is a simple implementation of Future that just wraps a constant value.

- map2 doesn't evaluate the call to f in a separate logical thread, in accord with our design choice of having fork be the sole function in the API for controlling parallelism.



Quiz

- ❖ Define $\text{map}[A,B](pa: \text{Par}[A])(f: A \Rightarrow B): \text{Par}[B]$ in terms of:
 - ❖ $\text{map2}[A,B,C] (a: \text{Par}[A], b: \text{Par}[B]) (f: (A,B) \Rightarrow C): \text{Par}[C]$
- ❖
- ❖ The fact that we can implement map in terms of map2 but not inversely, shows that map2 is strictly more powerful than map .
- ❖ This sort of thing happens a lot when we're designing libraries—often, a function that seems to be primitive will turn out to be expressible using some more powerful primitive.

Answer

- ❖ Define $\text{map}[A,B](\text{pa}: \text{Par}[A])(f: A \Rightarrow B): \text{Par}[B]$ in terms of:
 - ❖ $\text{map2}[A,B,C] (a: \text{Par}[A], b: \text{Par}[B]) (f: (A,B) \Rightarrow C): \text{Par}[C]$

❖

```
def map[A,B] (pa: Par[A]) (f: A => B) : Par[B] =  
  map2 (pa, unit ( ( ) ) ) ( (a,_) => f(a) )
```

Laws and Properties

- ❖ Consider a unit test (incidentally of the unit function):
 - ❖ `map(unit(1)) (_ + 1) == unit(2)`
- ❖ One can define “==” on the `Par[Int]` as follows:
 - ❖ `def equal[A] (e: ExecutorService) (p: Par[A], p2: Par[A]) : Boolean = p(e).get == p2(e).get`
- ❖ Such laws are useful as they can be turned into tests systematically

Conclusion

- ❖ Motivation for functional parallelism
- ❖ Designing the APIs for pure functional parallelism
- ❖ Re-used Java components